

Strings in C

A Programmer's Nightmare

Benjamin Brewster

Except as noted, all images copyrighted with Creative Commons licenses,
with attributions given whenever available

C Strings

- Strings in C are sequences of characters contiguously stored
 - Not a native type like `int` or `float` in more advanced languages
- A “string” terminates with the null character
`\0`
- That’s *it!* Any further programmatic use of strings requires functions and procedures that work within this format

Displaying C Strings with Formatted Printing

- Formatted means numbers correctly printed with text
- Formatted printing is done with:
 - `printf()` Prints to standard out
 - `sprintf()` Prints to a string (a char array)
 - `fprintf()` Prints to a file
- These functions look for null terminators to know when to stop

Basic C String Functions

Warning: do not use the == operator!

- Use the string library functions:
 - `strcmp()` Compares two strings for equality
 - `strlen()` Returns the length of the string in characters, not including null terminator
 - `strcpy()` Copies one string into another
 - `strcat()` Returns one string that is a concatenation of itself with another string
- n-character versions:
 - `strncpy()` Copy only n characters - won't null-terminate a full array, or actually prevent you from over-writing an array
 - `strncat()` Appends only a portion of a string to another

Declaring C Strings

- Three ways of declaring the same string

1. `char* mystring = "my string";`

2. `char mystring[] = "my string";`

3. `char mystring[10] = "my string";`

- Are they really the same? And why do we care in OS?



Declaring C Strings

- Three ways of declaring the same string

```
1. char* mystring      = "my string";
```

```
2. char  mystring[]    = "my string";
```

```
3. char  mystring[20]  = "my string";
```

- Are they really the same? And why do we care in OS?
- Because this one difference shows how close C is to the underlying memory management being performed by UNIX

Declaring C Strings

- Three ways of declaring the same string

- 1. `char* mystring = "my string";`

- 2. `char mystring[] = "my string";`

- 3. `char mystring[20] = "my string";`

- Are they really the same? And why do we care in OS?
- Because this one difference shows how close C is to the underlying memory management being performed by UNIX
- I.e. you need to know this, because otherwise you'll break all the things and not know why

Declaring C Strings – Method 1

- Three ways of declaring the same string

```
1. char* mystring = "my string";
```

```
2. char mystring[] = "my string";
```

```
3. char mystring[20] = "my string";
```

- At compile time, creates a sequence of bytes in the **read-only initialized data segment** portion of memory with the contents **"my string"**
- During execution, creates a pointer on the **stack** (automatic variable) called `mystring` that points to the read-only sequence of characters in the **data segment**
- `mystring` can be pointed to other addresses (it doesn't hold chars)

Declaring C Strings – Method 1 – Example

```
#include <stdio.h>
```

```
void main()  
{
```

```
    char* mystring = "my string";  
    printf("Var is: %s\n", mystring);  
    mystring[3] = 'Q';  
    printf("Var is: %s\n", mystring);
```

```
}
```

mystring is a pointer put on the stack

"my string" is a string literal defined and stored in the read-only portion of the data segment

Index 3 bytes off of where mystring is pointing too, then change whatever is there to 'Q'...

Result:

Declaring C Strings – Method 1 – Example

```
#include <stdio.h>
```

```
void main()  
{
```

```
    char* mystring = "my string";  
    printf("Var is: %s\n", mystring);  
    mystring[3] = 'Q';  
    printf("Var is: %s\n", mystring);
```

```
}
```

mystring is a pointer put on the stack

"my string" is a string literal defined and stored in the read-only portion of the data segment

Result:

```
Var is: my string
```

```
Segmentation fault (core dumped)
```

... except you can't do that, because your program cannot change memory in the read-only portion of the data segment

Declaring C Strings – Method 2

- Three ways of declaring the same string

```
1. char* mystring = "my string";
```

```
2. char mystring[] = "my string";
```

```
3. char mystring[20] = "my string";
```

- During execution, creates space for 10 bytes on the **stack** as an automatic variable, names that variable `mystring`
- Puts **"my string"** into the variable `mystring` with a null terminator after it
- The variable `mystring` is editable, as it is an array

Declaring C Strings – Method 2 – Example

```
#include <stdio.h>
```

mystring is an array



```
void main()  
{
```

```
    char mystring[] = "my string";  
    printf("Var is: %s\n", mystring);  
    mystring[3] = 'Q';  
    printf("Var is: %s\n", mystring);
```

```
}
```

Result:

```
Var is: my string
```

```
Var is: my Qtring
```

Declaring C Strings – Method 3

- Three ways of declaring the same string

```
1. char* mystring = "my string";
```

```
2. char mystring[] = "my string";
```

```
3. char mystring[20] = "my string";
```

- Creates space for 20 bytes on the **stack** as an automatic variable, names that variable `mystring`
- Puts **"my string"** into the variable `mystring` with a null terminator after it
- The variable `mystring` is editable, as it is an array

String Literals (Again)

- What's wrong with this code:

```
char* mystring = "my string";  
strcpy(mystring, "AA string");  
printf(mystring);
```



String Literals (Again)

- What's wrong with this code:

```
char* mystring = "my string";  
strcpy(mystring, "AA string");  
printf(mystring);
```

- String literals cannot be changed in C - they are initialized in the **read-only** section of the **initialized data segment**
- When is this error caught?
 - Only at run-time, as a seg-fault; this compiles just fine

Buffer Overrun

- What's wrong with this?

```
char fiveStr[5] = "five";  
strcpy(fiveStr, "five6");  
printf(fiveStr);
```



Buffer Overrun

- What's wrong with this?

```
char fiveStr[5] = "five";  
strcpy(fiveStr, "five6");  
printf(fiveStr);
```

- "five6" is too long to store in fiveStr
- When is this error caught?
 - Never!
 - Unless something you needed is overwritten and a segfault occurs because a just-accessed pointer no longer points to where it was supposed to!

Fully Initializing C String Arrays

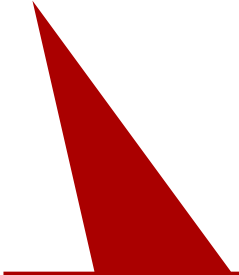
```
char mystring[20];  
strcpy(mystring, "my string");
```

m	y		s	t	r	i	n	g	\0	B	O	G	U	S	B	O	G	U	S
---	---	--	---	---	---	---	---	---	----	---	---	---	---	---	---	---	---	---	---

```
printf("%s", mystring);
```

Result:

my string



How do we deal with this uninitialized data?

What's In that Uninitialized Data?

```
$ cat cstring-array-unint.c
```

```
#include <stdio.h>
#include <string.h>
```

```
void main()
```

```
{
```

```
    int i = -5;
```

```
    char mystring[20];
```

```
    strcpy(mystring, "my string");
```

```
    printf("Char => Int :: ASCII Table Lookup\n");
```

```
    for (i = 0; i < 19; i++)
```

```
        printf("%c    => %d\n", mystring[i], mystring[i]);
```

```
}
```

```
$ gcc -o cstring-array-unint cstring-array-unint.c
```

```
$ ./cstring-array-unint
```

```
Char => Int :: ASCII Table Lookup
```

```
m    => 109
```

```
y    => 121
```

```
    => 32
```

```
s    => 115
```

```
t    => 116
```

```
r    => 114
```

```
i    => 105
```

```
n    => 110
```

```
g    => 103
```

```
    => 0
```

```
@    => 64
```

```
    => 0
```

```
    => 0
```

```
    => 0
```

```
    => 0
```

```
    => 0
```

```
@    => 64
```

```
W    => 87
```

```
█    => -26
```

Printing chars as
ints is a great way
to debug C string
arrays!

ASCII -26 summons Cthulu

Initializing C String Arrays

- The Bad

- Depending on how you declare them, C string arrays may be full of uninitialized data - it's best to clear them before use
- What happens if we somehow remove the automatic null terminator?

```
$ cat cstring-array.c
```

```
#include <stdio.h>
```

```
#include <string.h>
```

```
void main()
```

```
{
```

```
    int i = -5;
```

```
    char mystring[20];
```

```
    strcpy(mystring, "my string");
```

```
    printf("Var is: %s\n", mystring);
```

```
    mystring[3] = 'Q';
```

```
    printf("Var is: %s\n", mystring);
```

```
    mystring[9] = '#';
```

```
    mystring[19] = '\\0';
```

```
    for (i = 10; i < 19; i++)
```

```
        if (mystring[i] == '\\0')
```

```
            mystring[i] = '#';
```

```
    printf("Var is: %s\n", mystring);
```

```
}
```



Uninitialized!

Initializing C String Arrays

- The Bad

- Depending on how you declare them, C string arrays may be full of uninitialized data - it's best to clear them before use
- What happens if we somehow remove the automatic null terminator?

Different almost every time it runs, as memory is used

```
$ cat cstring-array.c
#include <stdio.h>
#include <string.h>

void main()
{
    int i = -5;
    char mystring[20];
    strcpy(mystring, "my string");
    printf("Var is: %s\n", mystring);
    mystring[3] = 'Q';
    printf("Var is: %s\n", mystring);
    mystring[9] = '#';
    mystring[19] = '\0';
    for (i = 10; i < 19; i++)
        if (mystring[i] == '\0')
            mystring[i] = '#';
    printf("Var is: %s\n", mystring);
}

$ gcc -o cstring-array cstring-array.c
$ ./cstring-array
Var is: my string
Var is: my Qtring
Var is: my Qtring#@####t
```

Uninitialized!

Initializing C String Arrays

- The Suspicious

- Depending on how you declare them, C string arrays may be full of uninitialized data - it's best to clear them before use

```
$ cat cstring-array.c
#include <stdio.h>
#include <string.h>

void main()
{
    int i = -5;
    char mystring[20] = "my string";
    printf("Var is: %s\n", mystring);
    mystring[3] = 'Q';
    printf("Var is: %s\n", mystring);
    mystring[9] = '#';
    mystring[19] = '\0';
    for (i = 10; i < 19; i++)
        if (mystring[i] == '\0')
            mystring[i] = '#';
    printf("Var is: %s\n", mystring);
}

$ gcc -o cstring-array cstring-array.c
$ ./cstring-array
Var is: my string
Var is: my Qtring
Var is: my Qtring#####
```

Seems to initialize
entire array to \0
but is this portable?

Initializing C String Arrays


- The Preferred

- Depending on how you declare them, C string arrays may be full of uninitialized data - it's best to clear them before use

```
$ cat cstring-array.c
#include <stdio.h>
#include <string.h>

void main()
{
    int i = -5;
    char mystring[20];
    memset(mystring, '\0', 20);
    strcpy(mystring, "my string");
    printf("Var is: %s\n", mystring);
    mystring[3] = 'Q';
    printf("Var is: %s\n", mystring);
    mystring[9] = '#';
    mystring[19] = '\0';
    for (i = 10; i < 19; i++)
        if (mystring[i] == '\0')
            mystring[i] = '#';
    printf("Var is: %s\n", mystring);
}

$ gcc -o cstring-array cstring-array.c
$ ./cstring-array
Var is: my string
Var is: my Qtring
Var is: my Qtring#####
```



Fully
Initialized

Meanwhile Back on the Ranch...

- C continues to provide dangerous string functions

strtok() :: String tokenizer

- Splits strings into chunks
- Makes your hair fall out
- Maxes out your credit cards
- Unfriends all your social media friends
- Sometimes the best/only tool for the job :/



strtok Example

```
char input[18] = "This.is my/string";
```

```
char* token = strtok(input, " ./");
```

This

```
token = strtok(NULL, " ./");
```

is

```
token = strtok(NULL, " ");
```

my/string

Changing the delimiter as strtok()
tokenizes the string is neat

A Major `strtok()` Drawback *(the first of many)*

```
char* input = "This.is my/string";  
char* token = strtok(input, " ./");  
token = strtok(NULL, " ./");  
token = strtok(NULL, " ");
```

- Fails miserably. Why?

A Major `strtok()` Drawback *(the first of many)*

```
char* input = "This.is my/string";  
char* token = strtok(input, " ./");  
token = strtok(NULL, " ./");  
token = strtok(NULL, " ");
```

- Fails miserably, crashing on execution: Why?
 - Because `input` is a string literal, and `strtok()` is about to mess with your strings

```
#include <stdio.h>
#include <string.h>

void main()
{
    char input[50];
    char* token = 0; // Set null pointer
    int inputsz = -5;
    int currChar = -5;

    memset(input, '\\0', 50);
    strcpy(input, "A.B C/D");
    inputsz = strlen(input);

    printf("Input: "); for (currChar = 0; currChar < inputsz; currChar++) printf("%2d ", input[currChar]);
    printf(" = \"%s\\", token: \"%s\\\"\\n", input, token);
    token = strtok(input, " ./");
    printf("Input: "); for (currChar = 0; currChar < inputsz; currChar++) printf("%2d ", input[currChar]);
    printf(" = \"%s\\", token: \"%s\\\"\\n", input, token);
    token = strtok(NULL, " ./");
    printf("Input: "); for (currChar = 0; currChar < inputsz; currChar++) printf("%2d ", input[currChar]);
    printf(" = \"%s\\", token: \"%s\\\"\\n", input, token);
    token = strtok(NULL, " ./");
    printf("Input: "); for (currChar = 0; currChar < inputsz; currChar++) printf("%2d ", input[currChar]);
    printf(" = \"%s\\", token: \"%s\\\"\\n", input, token);
}
```

strtok Example Results

Input: 65 46 66 32 67 47 68 = "A.B C/D", token: "(null) "

Input: 65 0 66 32 67 47 68 = "A", token: "A"

Input: 65 0 66 0 67 47 68 = "A", token: "B"

Input: 65 0 66 0 67 0 68 = "A", token: "C"

- input gets jacked up by strtok() as the delimiters encountered during parsing get nulled
- Further, this can only work because strtok() keeps a hidden static variable in the data segment up to date while parsing



Further strtok Horrors

- Not only does `strtok()` modify the input...
(You don't even specify which string to tokenize past the first call! Hidden vars!)

```
char input[18] = "This.is my/string";  
char* token = strtok(input, " ./");  
token = strtok(NULL, " ./");  
token = strtok(NULL, " ");
```

- Mixing calls of `strtok()` between different strings is not allowed because it can only process ONE string with its hidden variables!
 - But there is a `strtok_r()` that is re-entrant, allowing the mixing of calls

Horrors Explained

- This mixing of `strtok()` calls is easy to do on accident in a large program, especially with functions involved:

```
strtok(input1, ...)  
function()  
    strtok(input2, ...)  
strtok(input1, ...)
```

- The solution is to simply use a more modern language with a string type

Combining Declaration Methods

- What does this mean:

```
char* mystring[3];
```



Combining Declaration Methods

- What does this mean:

```
char* mystring[3];
```

Declare an array of pointers, each of which points to a string; each of these pointers can be pointed at either array names *or* string literals

Remember that an array name is a pointer to the first element's address in memory

Arrays of Pointers to Strings - Example

```
#include <stdio.h>
#include <string.h>

void main()
{
    int currElem = -5;
    int numElems = 3;
    char* mystring[numElems];
    char myarray[10];

    strcpy(myarray, "1ARRAY");

    printf("Size of char*: %d\n", sizeof(char*));
    printf("Size of one array element: %d\n", sizeof(mystring[0]));
    printf("Size of all array elements: %d\n", sizeof(mystring));
    printf("Number of elements in array: %d = %d\n", sizeof(mystring) / sizeof(mystring[0]), numElems);

    //strcpy(mystring[0], "strcpy string");           // Causes seg fault, that's a pointer!
    //printf("mystring[0]: %s\n", mystring[0]);
    mystring[0] = "string literal";                   // Set the first pointer to point to the address of a string literal
    printf("mystring[0]: %s\n", mystring[0]);         // (which is the address of the literal's first element)
    mystring[0] = myarray;                           // Set the first pointer to point to the name of a C string array
    printf("mystring[0]: %s\n", mystring[0]);         // (which is the address of the array's first element)
}
```

Results:

Size of char*: 8
Size of one array element: 8
Size of all array elements: 24
Number of elements in array: 3 = 3
mystring[0]: string literal
mystring[0]: 1ARRAY

Combining Declaration Methods

- What does this mean:

```
char* mystring[3];
```

Declare an array of pointers, each of which points to a string; each of these pointers can be pointed at either array names *or* string literals

Remember that an array name is a pointer to the first element's address in memory

We can change where the pointers point, but how do we create new strings for this new array to hold?

Dynamically Allocating a String

- To create a string variable dynamically, and thus use it like an array, use `malloc()` and `free()`:

Note that `char* mystring` is editable!

```
$ gcc -o malloctest malloctest.c
$ malloctest
yay! literal
yayQ literal
```

```
#include <stdlib.h>
#include <stdio.h>
#include <string.h>

void main()
{
    char* mystring;
    char* literal = "literal";

    mystring = malloc(20 * sizeof(char));

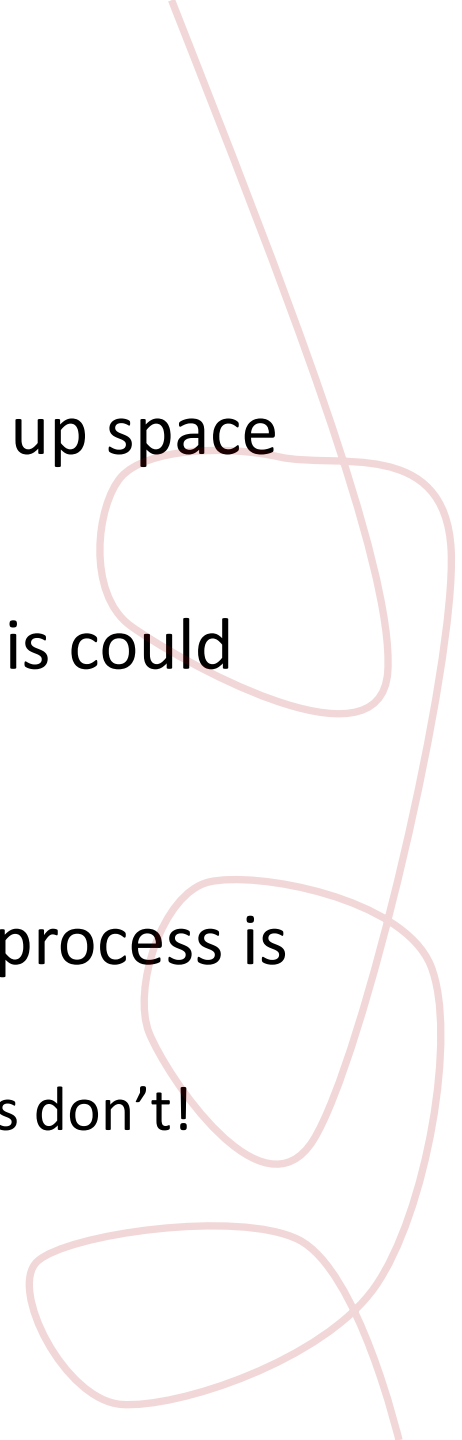
    if (mystring == 0)
        printf("malloc() failed!\n");

    memset(mystring, '\\0', 20);

    sprintf(mystring, "yay! %s\\n", literal);
    printf("%s", mystring);
    mystring[3] = 'Q';
    printf("%s", mystring);

    free(mystring);
}
```

Malloc Memory Leaks

- If you don't free dynamically allocated memory, it still takes up space
 - If you have a long-running program, like a server process, this could eventually use up all of your memory
 - Process memory is normally all freed automatically when a process is terminated
 - At least in UNIX, Windows, etc. - some real-time operating systems don't!
- 

Malloc Memory Leaks

- Here's a classic way to hide and cause a leak:

```
char* mystring = malloc(20 * sizeof(char));  
...  
mystring = "hello";
```

- This leaks because you no longer have the start address of the dynamically allocated space; `mystring` now points to a string literal

```
free(mystring); // And if you try this later, it fails spectacularly
```

Spectacular Failing



Same program, but let's just
put this right in here...

```
#include <stdlib.h>
#include <stdio.h>
#include <string.h>

void main()
{
    char* mystring;
    char* literal = "literal";

    mystring = malloc(20 * sizeof(char));

    if (mystring == 0)
        printf("malloc() failed!\n");

    memset(mystring, '\0', 20);

    sprintf(mystring, "yay! %s\n", literal);
    printf("%s", mystring);
    mystring[3] = 'Q';
    printf("%s", mystring);
    mystring = "test\n";
    free(mystring);
}
```

Spectacular Failing Results

```
$ malloctest
yay! literal
yayQ literal
*** Error in `malloctest': free(): invalid pointer: 0x000000000400805 ***
===== Backtrace: =====
/lib64/libc.so.6(+0x7d053) [0x7fc92849c053]
malloctest[0x40074a]
/lib64/libc.so.6(__libc_start_main+0xf5) [0x7fc928440b15]
malloctest[0x4005d9]
===== Memory map: =====
00400000-00401000 r-xp 00000000 00:39 3238103636 /nfs/stak/faculty/b/brewsteb/tempdir/malloctest
00600000-00601000 r--p 00000000 00:39 3238103636 /nfs/stak/faculty/b/brewsteb/tempdir/malloctest
00601000-00602000 rw-p 00001000 00:39 3238103636 /nfs/stak/faculty/b/brewsteb/tempdir/malloctest
01195000-011b6000 rw-p 00000000 00:00 0 [heap]
7fc924000000-7fc924021000 rw-p 00000000 00:00 0
7fc924021000-7fc928000000 ---p 00000000 00:00 0
7fc928209000-7fc92821e000 r-xp 00000000 fd:02 16777347 /usr/lib64/libgcc_s-4.8.5-20150702.so.1
7fc92821e000-7fc92841d000 ---p 00015000 fd:02 16777347 /usr/lib64/libgcc_s-4.8.5-20150702.so.1
7fc92841d000-7fc92841e000 r--p 00014000 fd:02 16777347 /usr/lib64/libgcc_s-4.8.5-20150702.so.1
7fc92841e000-7fc92841f000 rw-p 00015000 fd:02 16777347 /usr/lib64/libgcc_s-4.8.5-20150702.so.1
7fc92841f000-7fc9285d6000 r-xp 00000000 fd:02 16811513 /usr/lib64/libc-2.17.so
7fc9285d6000-7fc9287d6000 ---p 001b7000 fd:02 16811513 /usr/lib64/libc-2.17.so
7fc9287d6000-7fc9287da000 r--p 001b7000 fd:02 16811513 /usr/lib64/libc-2.17.so
7fc9287da000-7fc9287dc000 rw-p 001bb000 fd:02 16811513 /usr/lib64/libc-2.17.so
7fc9287dc000-7fc9287e1000 rw-p 00000000 00:00 0
7fc9287e1000-7fc928802000 r-xp 00000000 fd:02 16811685 /usr/lib64/ld-2.17.so
7fc9289cf000-7fc9289d2000 rw-p 00000000 00:00 0
7fc9289ff000-7fc928a02000 rw-p 00000000 00:00 0
7fc928a02000-7fc928a03000 r--p 00021000 fd:02 16811685 /usr/lib64/ld-2.17.so
7fc928a03000-7fc928a04000 rw-p 00022000 fd:02 16811685 /usr/lib64/ld-2.17.so
7fc928a04000-7fc928a05000 rw-p 00000000 00:00 0
7ffe32184000-7ffe321a5000 rw-p 00000000 00:00 0
7ffe321cd000-7ffe321cf000 r-xp 00000000 00:00 0
fffffffffff60000-fffffffffff601000 r-xp 00000000 00:00 0
Aborted (core dumped)
```

